

Previously...

1

- Introduction and history of AI
- Agents: rational agent, agent types
- Task environment (i.e., PEAS) and its properties

Today...

- Uninformed Search

SOLVING PROBLEMS BY SEARCHING

ALMA Chapter 3.1 – 3.4

Outline

3

- Problem-solving agent: A type of goal-based agent using atomic representation
- Problem definition/formulation
- Example problems
- Uninformed search algorithms

Problem-Solving Agent

4

- Assume that task environment is fully observable, deterministic, and discrete → Solution is fixed sequence of actions

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  persistent: seq, an action sequence, initially empty
               state, some description of the current world state
               goal, a goal, initially null
               problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
    if seq = failure then return a null action
  action ← FIRST(seq)
  seq ← REST(seq)
  return action
```

- Are percepts needed to execute rest of sequence of actions?
No! → Open-loop system

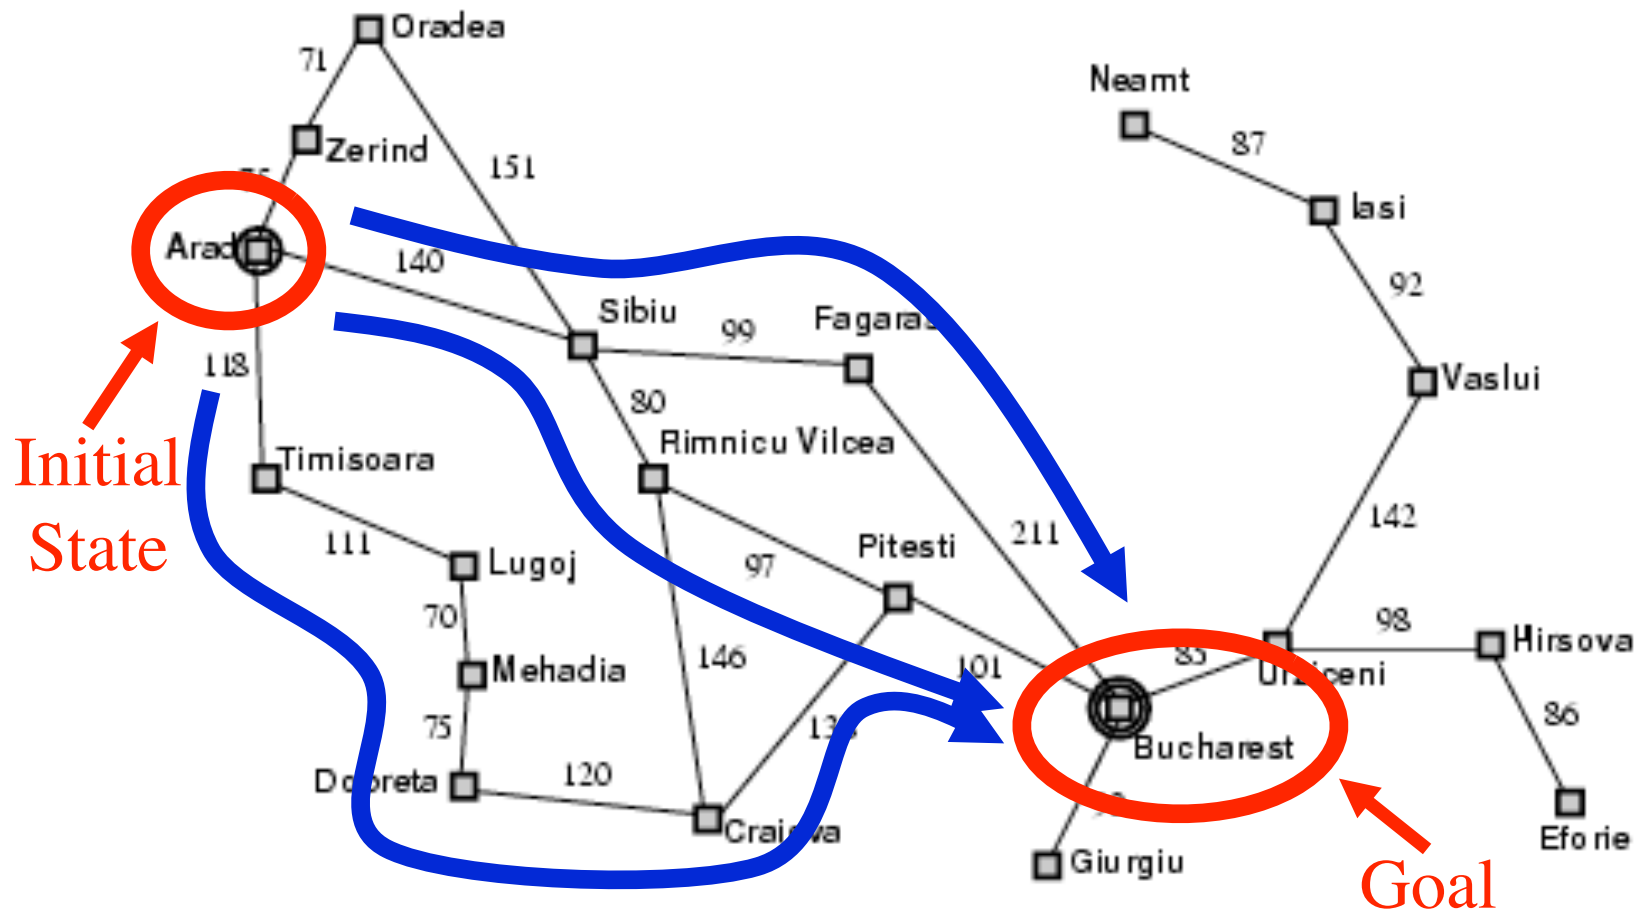
Example: Romania

5

- On tour in Romania; currently in Arad
- Flight leaves tomorrow from Bucharest
- Formulate goal:
 - ▣ be in Bucharest
- Formulate problem:
 - ▣ **states**: various cities
 - ▣ **actions**: drive between cities
- Solve search problem:
 - ▣ sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

Example: Romania

6



Problem Formulation

7

A problem is defined by 6 components:

1. Initial state : e.g., $In(Arad)$
2. States : e.g., $\{In(Arad), In(Sibiu), In(Timisoara), In(Zerind), \dots\}$
3. Actions : $ACTIONS(s)$ returns set of actions that can be executed in state s .
e.g., $ACTIONS(In(Arad)) = \{Go(Sibiu), Go(Timisoara), Go(Zerind)\}$
4. Transition model : $RESULT(s, a)$ returns state that results from doing action a in state s . e.g., $RESULT(In(Arad), Go(Zerind)) = In(Zerind)$

Problem Formulation

8

A problem is defined by 6 components:

5. Goal test determines whether a given state is a goal state
 - Explicit set of goal states, e.g., $\{In(Bucharest)\}$
 - Implicit function, e.g., $Checkmate(s)$
6. Path cost is a cost function assigning a numeric cost to each path
 - Additive. e.g., sum of distances, number of actions executed
 - $c(s, a, s')$ is the **step cost** of taking action a in state s to reach state s' .
Assumed to be ≥ 0

A **solution** is a sequence of actions leading from the initial state to a goal state


Defining State Space

9

- Real world is absurdly complex → state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
e.g., When in state *In(Arad)*, action *Go(Zerind)* represents a complex set of possible routes, detours, rest stops, etc.
- (Abstract) solution = set of real paths that are solutions in the real world
- **Valid** abstraction: e.g., for **any** real state “in Arad”, there is a real path to **some** real state “in Zerind”
- **Useful** abstraction: executing each abstract action is “easier” than the original problem

Example: Vacuum World

10



(LEFT) State-space graph

1. Initial state : Any state
2. States : Unique integer to represent agent and dirt locations
3. Actions : $\{Left, Right, Suck\}$

4. Transition model :

4. Goal test : no dirt at all locations

5. Path cost : number of steps in path (step cost of 1)

Example: 8-Puzzle

11

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

1. Initial state : Any state
2. States : Locations of tiles
3. Actions : Move blank *Left*, *Right*, *Up*, or *Down*

4. Transition model : e.g., If we apply *Left* to start state, the resulting state has 5 and blank switched
5. Goal test : goal state is given
6. Path cost : number of steps in path (step cost of 1)

Note: Optimal solution of n-puzzle family is NP-hard.

Tree Search Algorithms

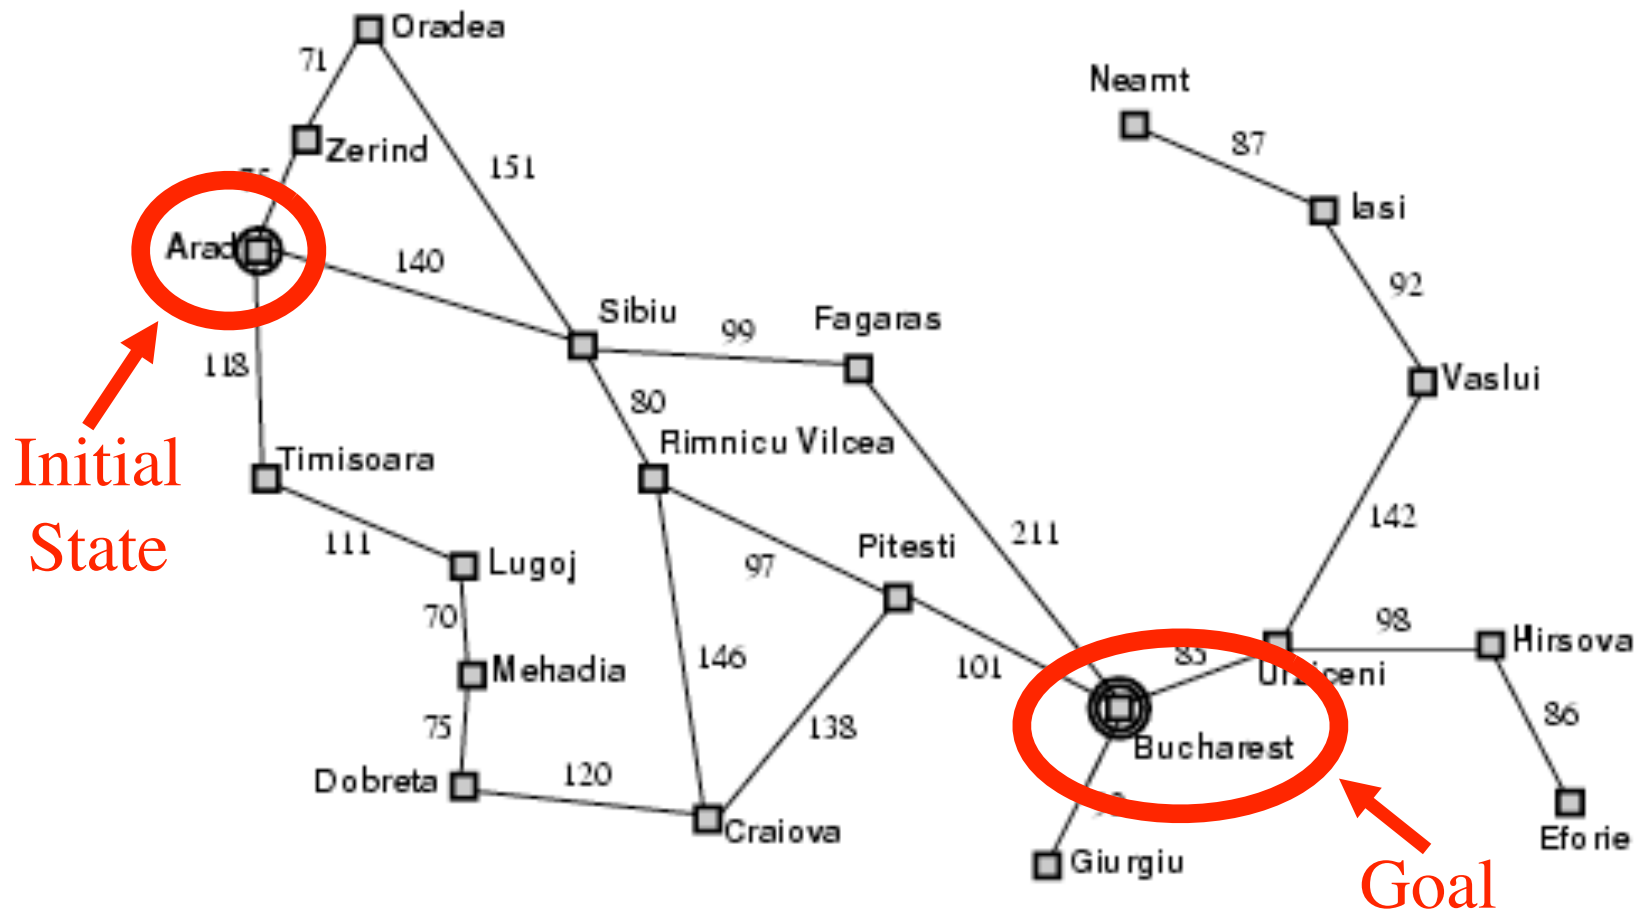
12

Basic idea: offline, simulated exploration of state space by **generating** successors of already-explored states (a.k.a. **expanding** states)

```
function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier
```

Example: Romania

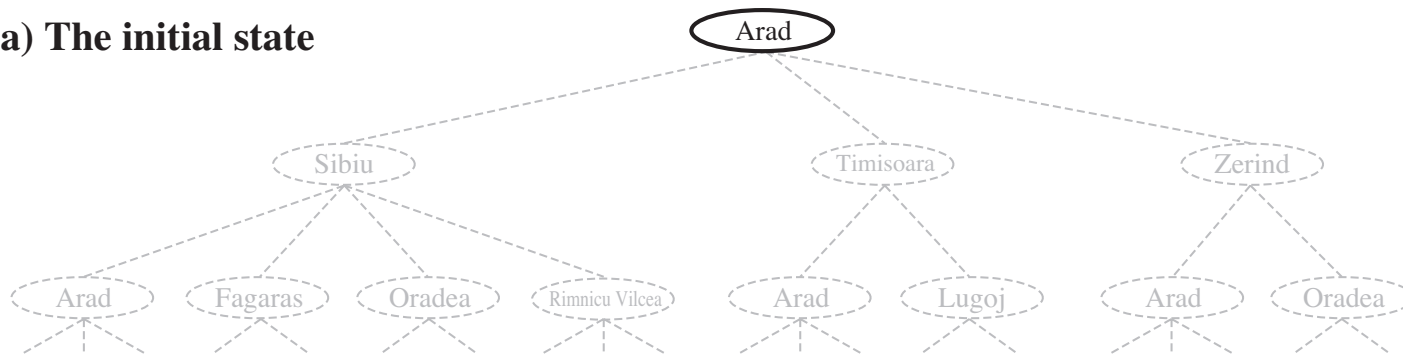
13



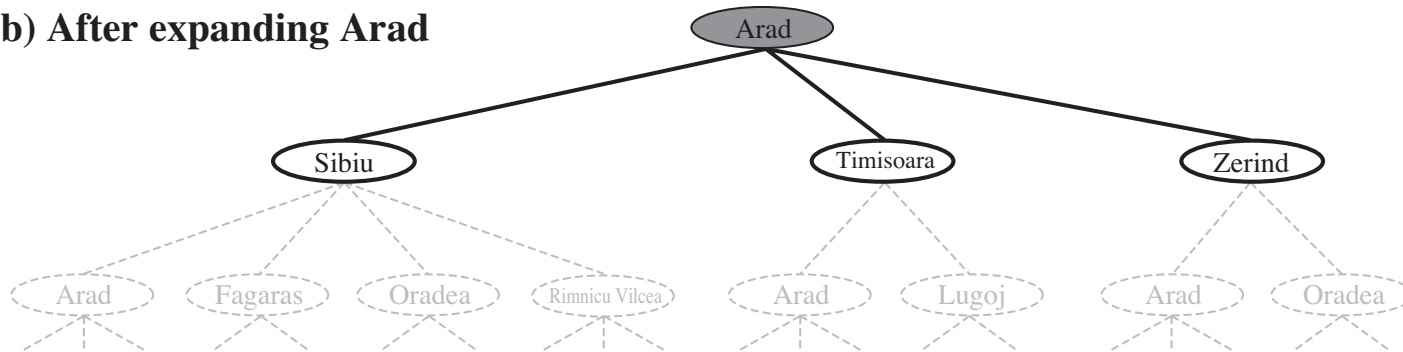
Tree Search Example

14

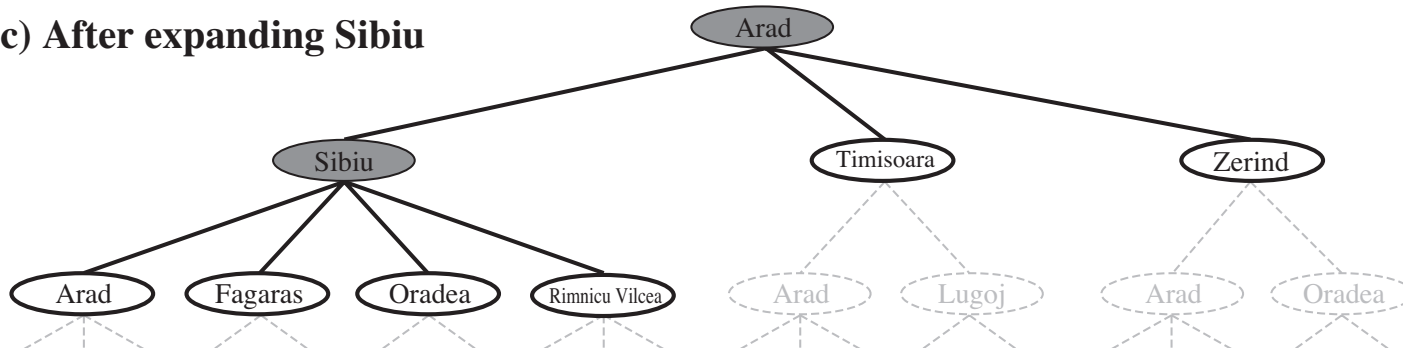
(a) The initial state



(b) After expanding Arad




(c) After expanding Sibiu



Implementation: States vs. Nodes

15

- A **state** represents a physical configuration
- A **node** is a data structure constituting part of search tree. It includes **state**, **parent node**, **action**, and **path cost** $g(n)$.
- Two different nodes are allowed to contain same world state



Search Strategies

16

A search strategy is defined by picking the **order of node expansion**.

- Strategies are evaluated with the following criteria:
 - **completeness**: does it always find a solution if one exists?
 - **optimality**: does it always find a least-cost solution?
 - **time complexity**: number of nodes generated during search
 - **space complexity**: maximum number of nodes in memory
- Time and space complexity are measured in terms of
 - b : maximum branching factor of search tree
 - d : depth of shallowest goal node
 - m : maximum depth of search tree (may be ∞)

Uninformed Search Strategies

17

- **Uninformed** search strategies use only the information available in the problem definition
- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search

Breadth-First Search (BFS)


18

- **Idea:** Expand shallowest unexpanded node
- **Implementation:** *Frontier* is a FIFO queue, i.e., insert new successors at the end

Breadth-First Search (BFS)

19


- **Idea:** Expand shallowest unexpanded node



Breadth-First Search (BFS)

20


- **Idea:** Expand shallowest unexpanded node



Breadth-First Search (BFS)

21


- **Idea:** Expand shallowest unexpanded node



Breadth-First Search (BFS)

22

- **Idea:** Expand shallowest unexpanded node



Properties of Breadth-First Search

23

- Complete? Yes (if b is finite)
- Optimal? Yes (if step cost = 1)
- Time? $b + b^2 + b^3 + \dots + b^d = O(b^d)$
- Space? $O(b^d)$ (keeps frontier nodes in memory)

Space is the bigger problem (more than time)

Uniform-Cost Search (UCS)

24

- Idea: Expand least-path-cost unexpanded node
- Implementation:
Frontier = priority queue ordered by path cost g
- Similar to BFS if all step costs are equal
- Complete? Yes, if step cost $\geq \epsilon$
- Optimal? Yes – nodes expanded in increasing order of path cost $g(n)$
- Time? # of nodes with (path cost $g \leq \text{cost } C^*$ of optimal solution):
 $O(b^{1+\text{floor}(C^*/\epsilon)})$
- Space? # of nodes with (path cost $g \leq \text{cost } C^*$ of optimal solution):
 $O(b^{1+\text{floor}(C^*/\epsilon)})$

Depth-First Search


25

- Idea: Expand deepest unexpanded node
- Implementation: *Frontier* = LIFO stack, i.e., insert successors at the front

Depth-First Search

26

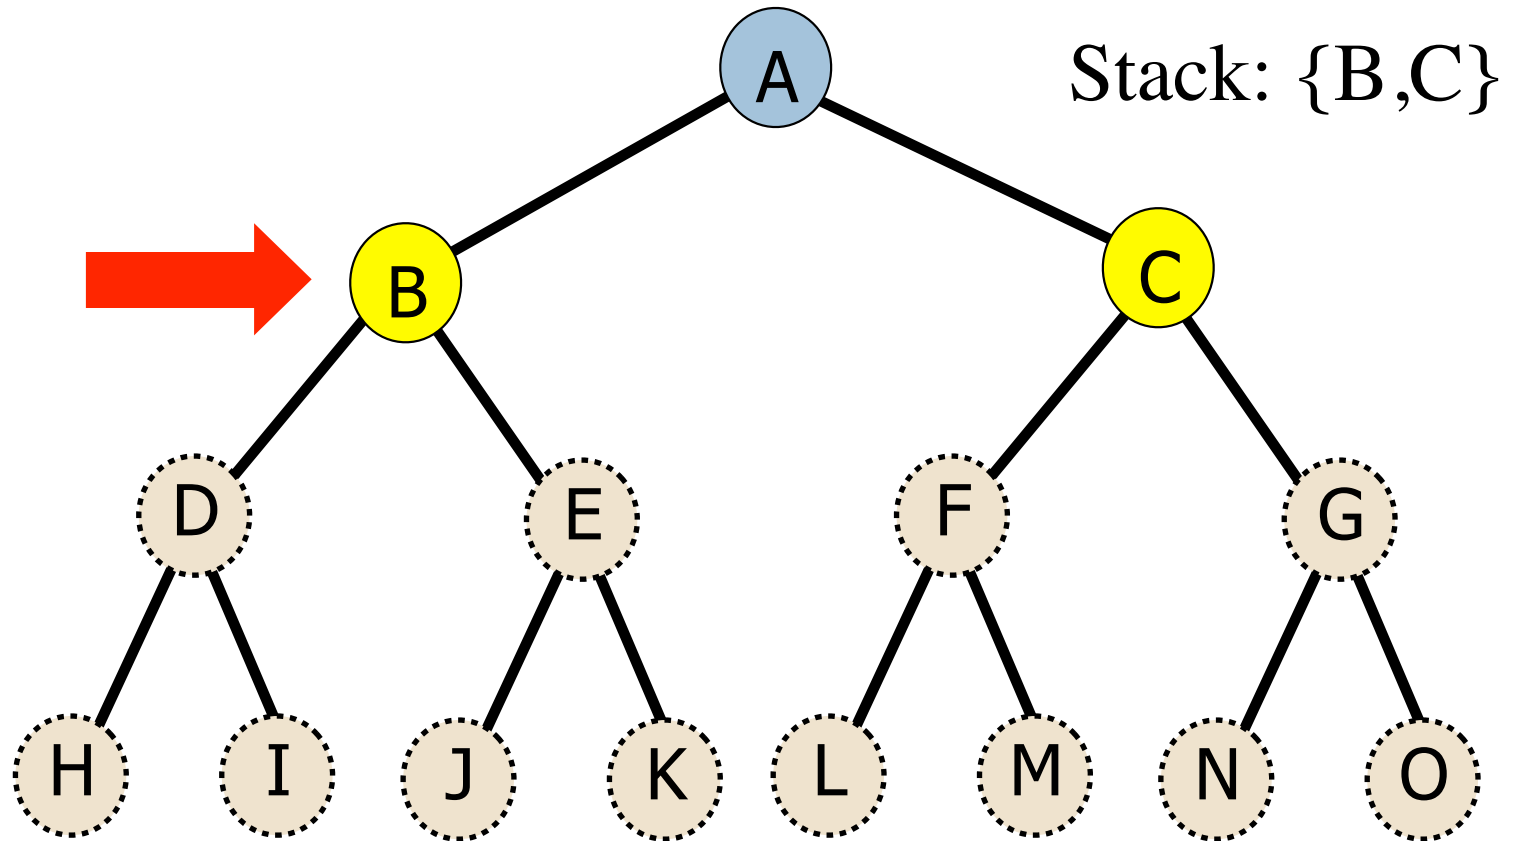
- Idea: Expand deepest unexpanded node



Depth-First Search

27

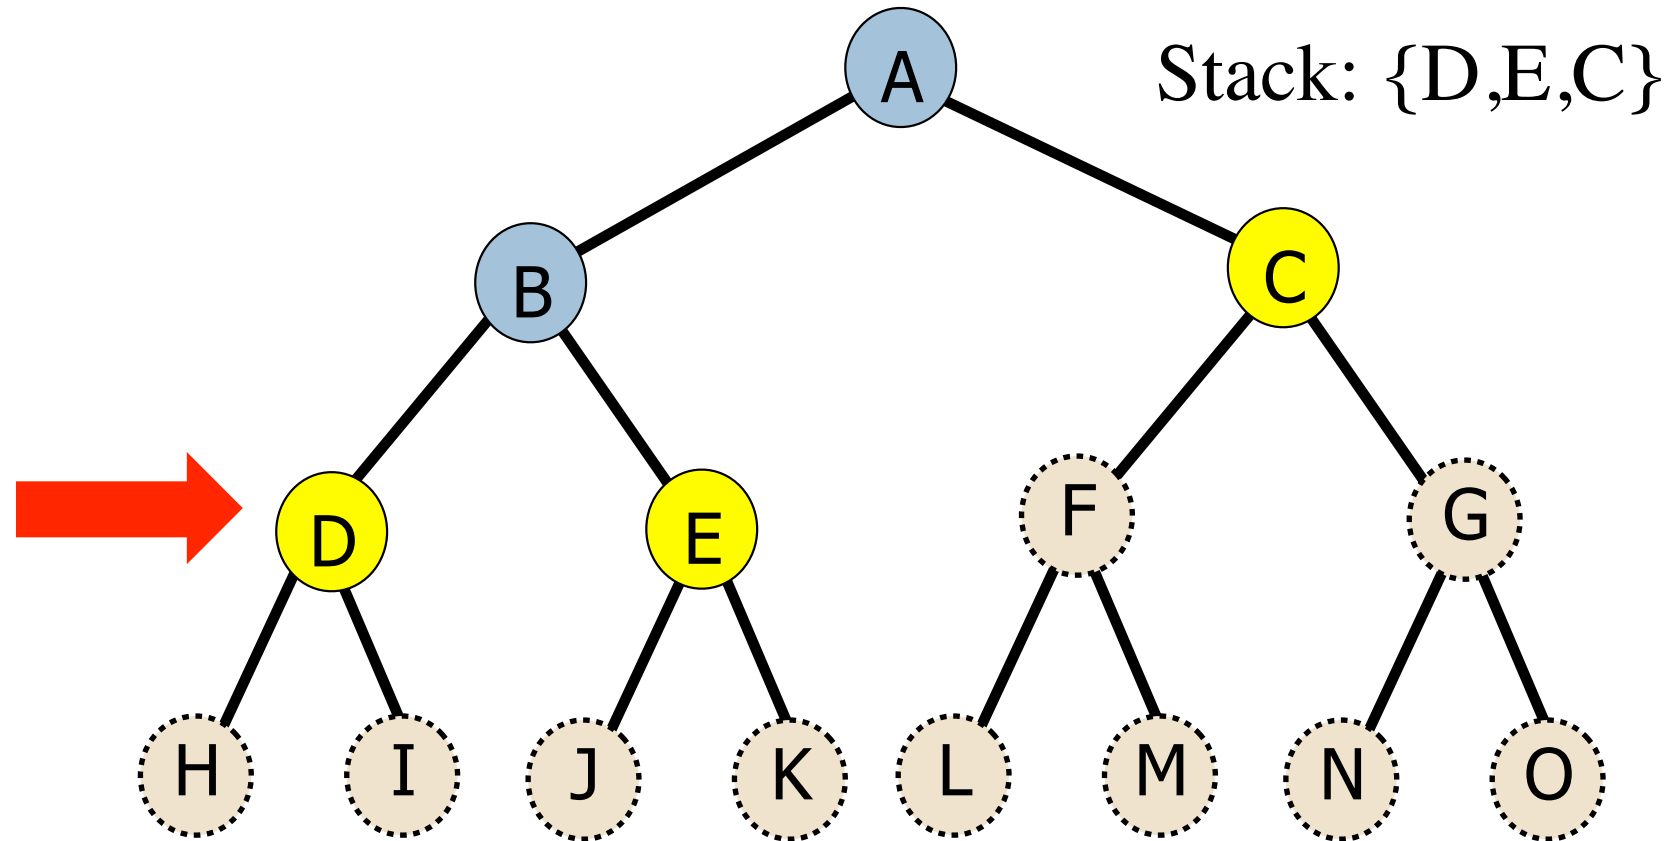
- Idea: Expand deepest unexpanded node



Depth-First Search

28

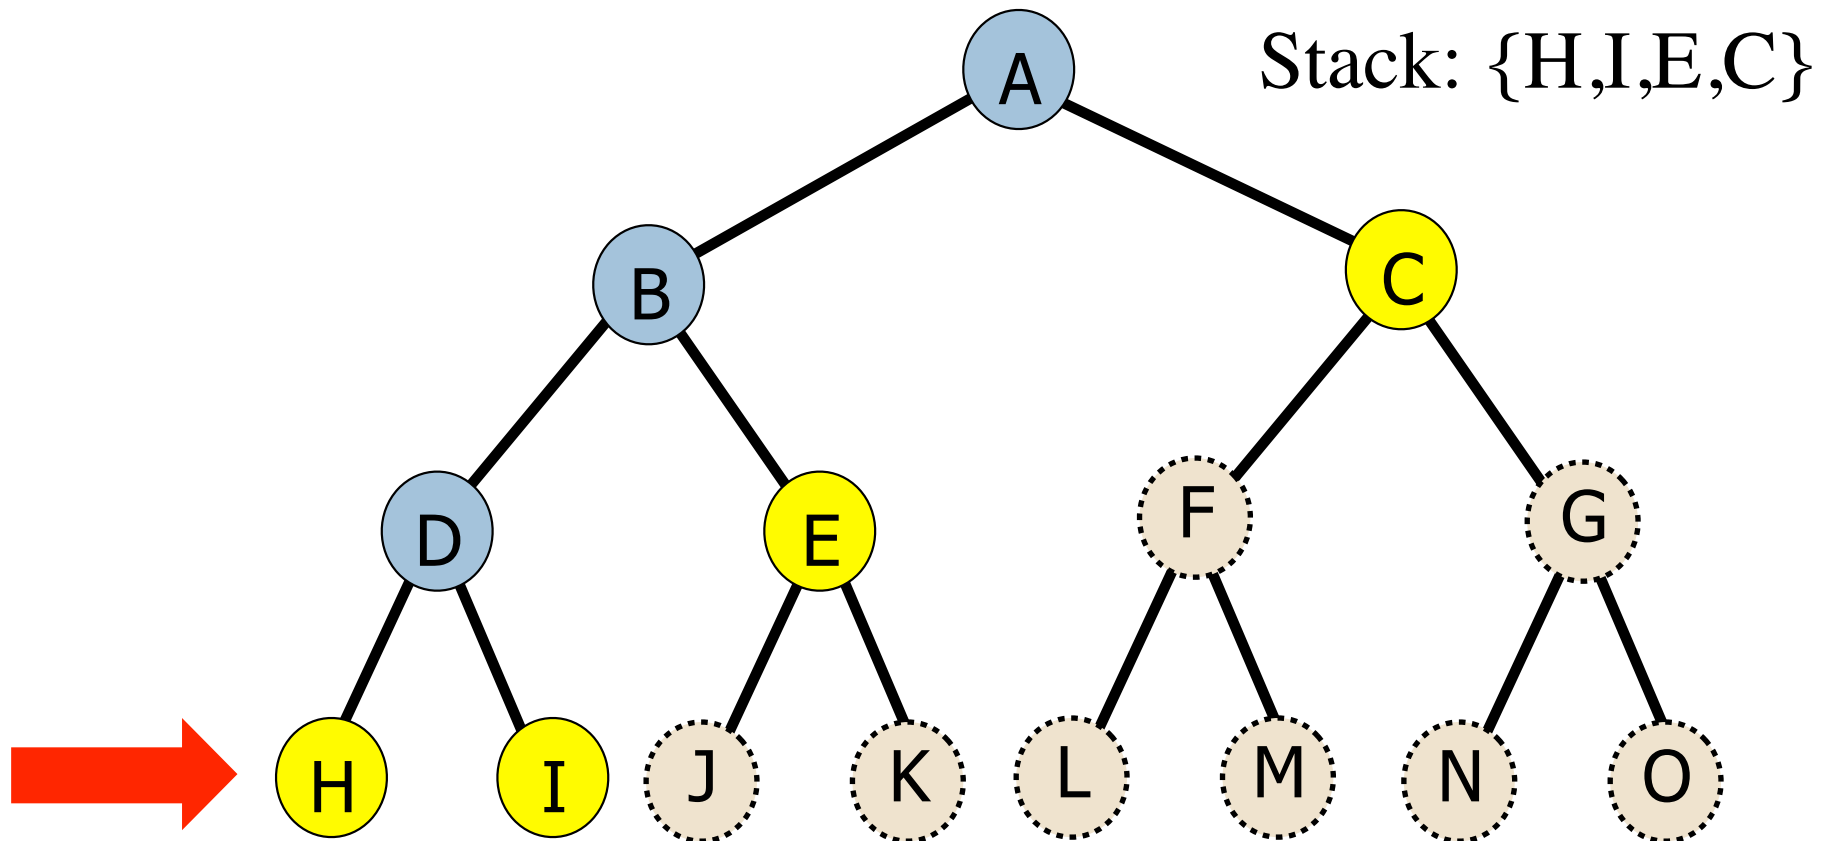
- Idea: Expand deepest unexpanded node



Depth-First Search

29


- Idea: Expand deepest unexpanded node



Depth-First Search

30

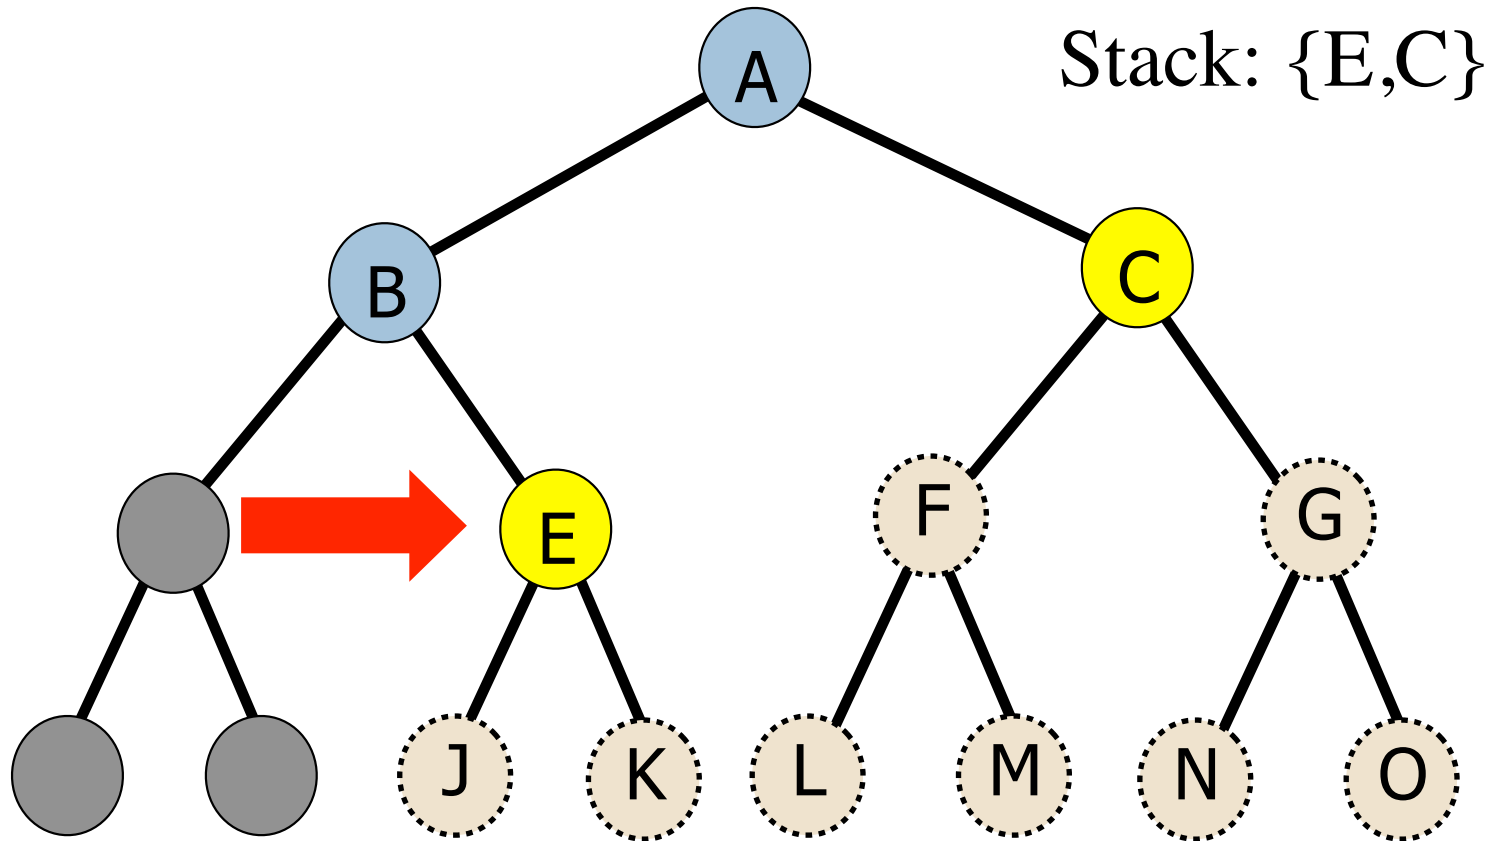
- Idea: Expand deepest unexpanded node



Depth-First Search

31


- Idea: Expand deepest unexpanded node



Depth-First Search

32


- Idea: Expand deepest unexpanded node



Depth-First Search

33


- Idea: Expand deepest unexpanded node



Depth-First Search

34


- Idea: Expand deepest unexpanded node



Depth-First Search

35


- Idea: Expand deepest unexpanded node



Depth-First Search

36


- Idea: Expand deepest unexpanded node



Depth-First Search

37


- Idea: Expand deepest unexpanded node



Depth-First Search

38


- Idea: Expand deepest unexpanded node



Depth-First Search

39


- Idea: Expand deepest unexpanded node



Depth-First Search

40

- Idea: Expand deepest unexpanded node



Properties of Depth-First Search

41

- Complete?

No: fails in infinite-depth search tree.

Can be modified to avoid repeated states along path → complete in finite state space

- Optimal?

No.

- Time?

$O(b^m)$: terrible if m is much larger than d .

But, if solutions are dense, it may be much faster than BFS.

- Space?

$O(bm)$, i.e., linear space!

Depth-Limited Search (DLS)

42

- Idea: Depth-first search with depth limit l , i.e., nodes at depth l have no successors
- Recursive implementation:

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns a solution, or failure/cutoff  
  return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)  
  
function RECURSIVE-DLS(node, problem, limit) returns a solution, or failure/cutoff  
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)  
  else if limit = 0 then return cutoff  
  else  
    cutoff_occurred?  $\leftarrow$  false  
    for each action in problem.ACTIONS(node.STATE) do  
      child  $\leftarrow$  CHILD-NODE(problem, node, action)  
      result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit - 1)  
      if result = cutoff then cutoff_occurred?  $\leftarrow$  true  
      else if result  $\neq$  failure then return result  
  if cutoff_occurred? then return cutoff else return failure
```

Iterative Deepening Search (IDS)

43

- Idea: Performs DLSs with increasing depth limit until goal node is found
- IDS is preferred if state space is large and depth of solution is unknown
- Implementation:

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
```

Iterative Deepening Search ($l = 0$)

44




Limit = 0



Iterative Deepening Search ($l = 1$)

45


Limit = 1



Iterative Deepening Search ($l = 2$)

46


Limit = 2



Iterative Deepening Search ($l = 3$)

47

Limit = 3



Iterative Deepening Search (IDS)

48

- Number of nodes generated in DLS (or BFS) to depth d with branching factor b :

$$N_{\text{DLS}} = b^0 + b^1 + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$

- Number of nodes generated in IDS to depth d with branching factor b :

$$N_{\text{IDS}} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

- For $b = 10$, $d = 5$,

- $N_{\text{DLS}} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$

- $N_{\text{IDS}} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$

- Overhead = $(123,456 - 111,111)/111,111 = 11\%$

Properties of Iterative Deepening Search

49

- Complete? Yes (if b is finite)
- Optimal? Yes (if step cost = 1)
- Time? $(d+1)b^0 + db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space? $O(bd)$

Comparing Tree Search Strategies

50



Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

1. BFS and IDS are complete if b is finite.
2. UCS is complete if b is finite and step cost $\geq \epsilon$
3. BFS and IDS are optimal if step costs are identical.

Repeated States

51

- Failure to detect repeated states can turn a linear problem into an exponential one!



Dealing with Repeated States

52

- How do we avoid repeats/loops?
Remember the nodes that are already visited!

Graph Search Algorithms



53

```
function GRAPH-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  initialize the explored set to be empty
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    add the node to the explored set
    expand the chosen node, adding the resulting nodes to the frontier
    only if not in the frontier or explored set
```

Bidirectional Search

54

- Simultaneously search both forward (from the initial state) and backward (from the goal state)
- Goal test: frontier of the two searches meet
- Intuition: $2 \times O(b^{d/2})$ is smaller than $O(b^d)$



Bidirectional Search

55

- Numerical Example ($b = 10, d = 6$)
 - Bi-directional search finds solution at $d = 3$ for both forward and backward search. Assuming BFS in each half, 2222 nodes are generated (as compared to 1,111,110 for full BFS).
- Implementation issues:
 - Actions are reversible: predecessor = successor
 - There may be many possible goal states
 - Construct a dummy goal state whose immediate predecessors are all the actual goal states
 - Check if a node appears in the “other” search tree
 - Using different search strategies for each half.