

Previously...

1

- **Uninformed** search strategies use only the information available in the problem definition
 - ▣ Breadth-first search
 - ▣ Uniform-cost search
 - ▣ Depth-first search
 - ▣ Depth-limited search
 - ▣ Iterative deepening search
- Which is the best search strategy?

Comparing Tree Search Strategies

2

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

1. BFS and IDS are complete if b is finite.
2. UCS is complete if b is finite and step cost $\geq \epsilon$
3. BFS and IDS are optimal if step costs are identical.

Choosing a Search Strategy

3

- Depends on the problem:
 - ▣ Finite/infinite depth of search tree
 - ▣ Known/unknown solution depth
 - ▣ Repeated states
 - ▣ Identical/non-identical step costs
 - ▣ Completeness and optimality needed?
 - ▣ Resource constraints (e.g., time, space)

Faster Search?

4

- Yes! Represent problem-specific knowledge as heuristics to guide search
- Today:
 - ▣ Informed (heuristic) search
 - ▣ Order of node expansion still matters: Which nodes are “more promising”?

INFORMED SEARCH

ALMA Chapter 3.5.1 – 3.5.2, 3.6.1 – 3.6.2


Outline

6

- Best-first search
- Greedy best-first search
- A^* search
- Heuristics

CMU's BOSS: DARPA Urban Challenge 2007

7



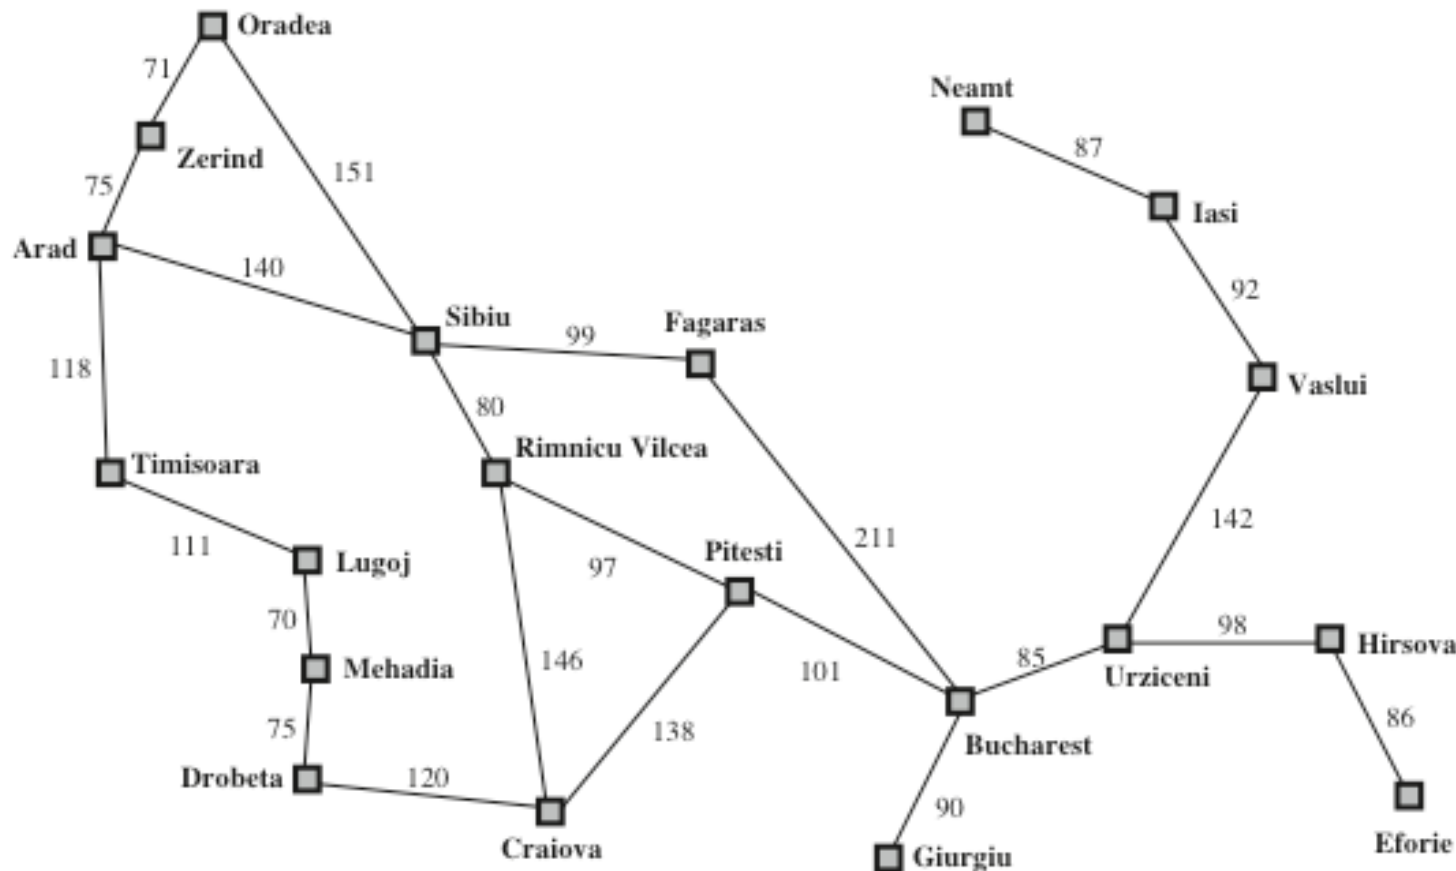
Best-First Search

8

- Idea: use an **evaluation function** $f(n)$ for each node n
 - ▣ Cost estimate
 - ▣ Expand node with lowest evaluation/cost first
- Implementation:
Frontier = priority queue ordered by nondecreasing cost f
- Special cases (different choices of f):
 - ▣ Greedy best-first search
 - ▣ A* search

Romania with Step Costs (km)

9



Values of h_{SLD} - straight-line distances to Bucharest

Arad	366
Bucharest	0
Craiova	160
Drobeta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	100
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy Best-First Search

10

- Evaluation function $f(n)$
= $h(n)$ (**h**euristic function)
= estimated cost of cheapest path from n to *goal*
- e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal

Greedy Best-First Search Example

11



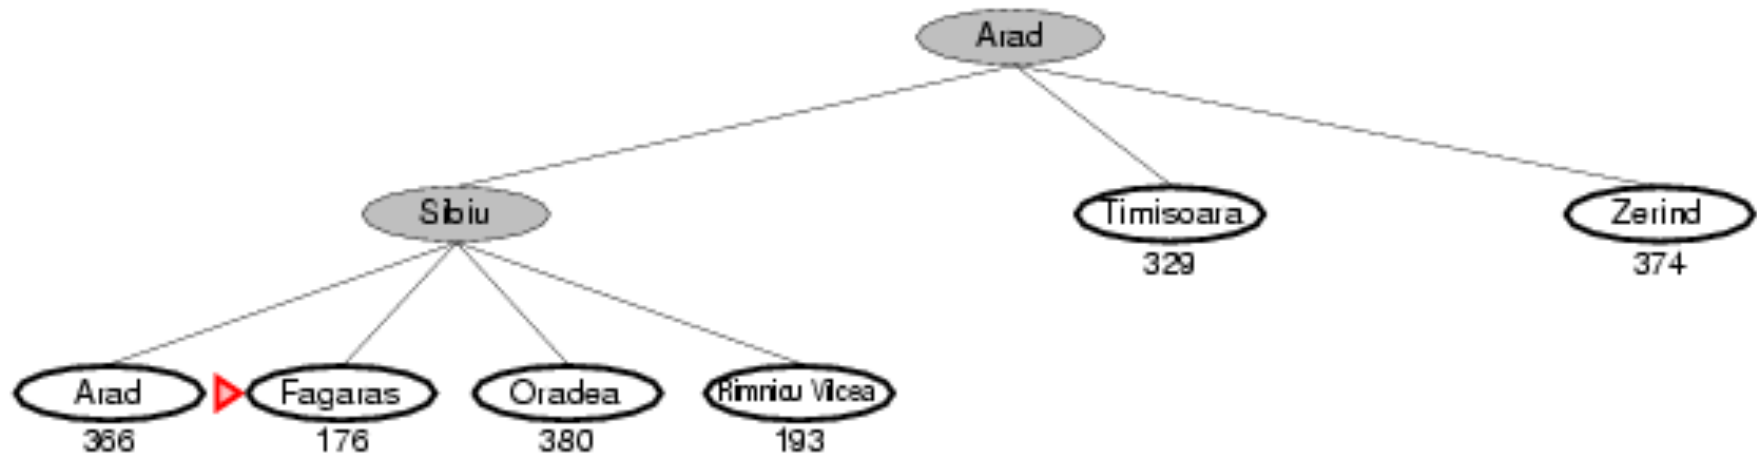
Greedy Best-First Search Example

12



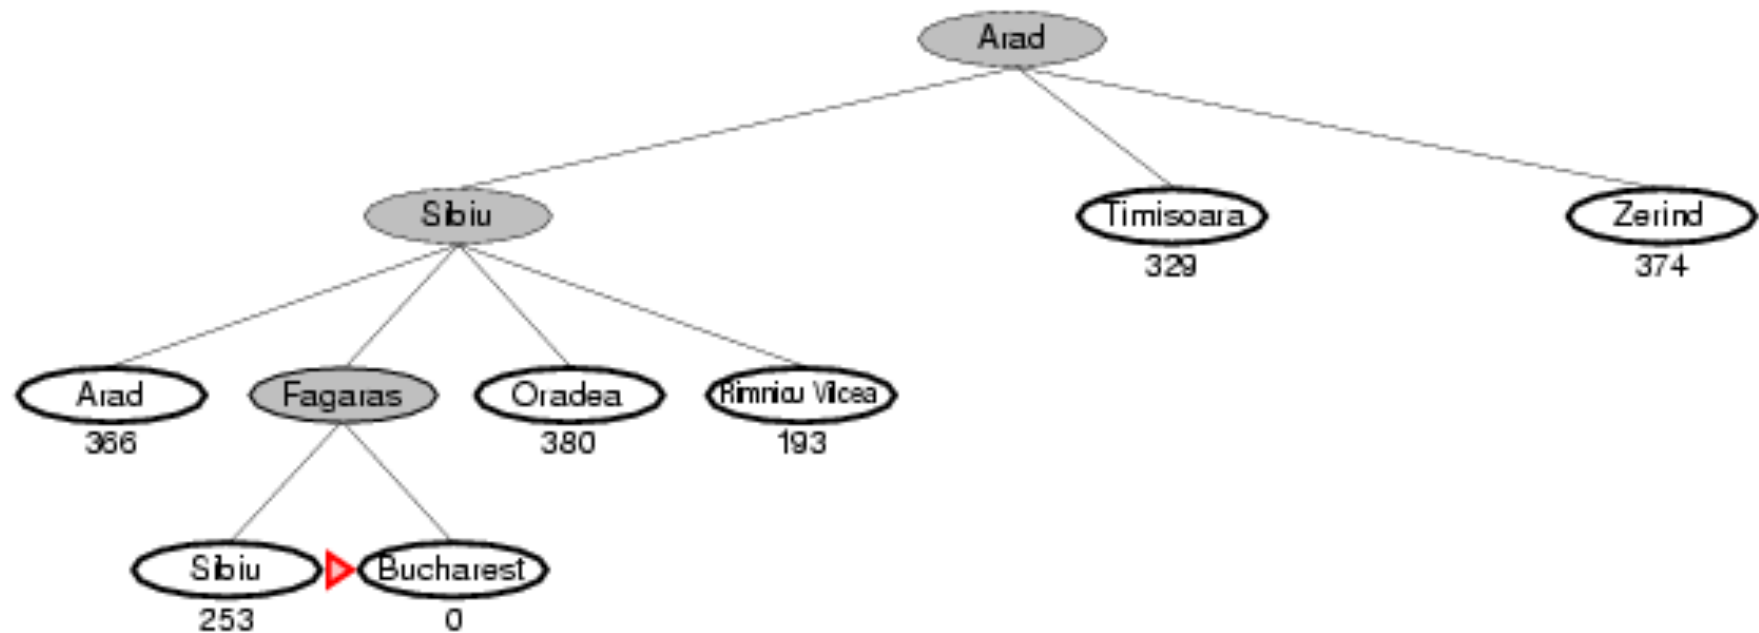
Greedy Best-First Search Example

13



Greedy Best-First Search Example

14



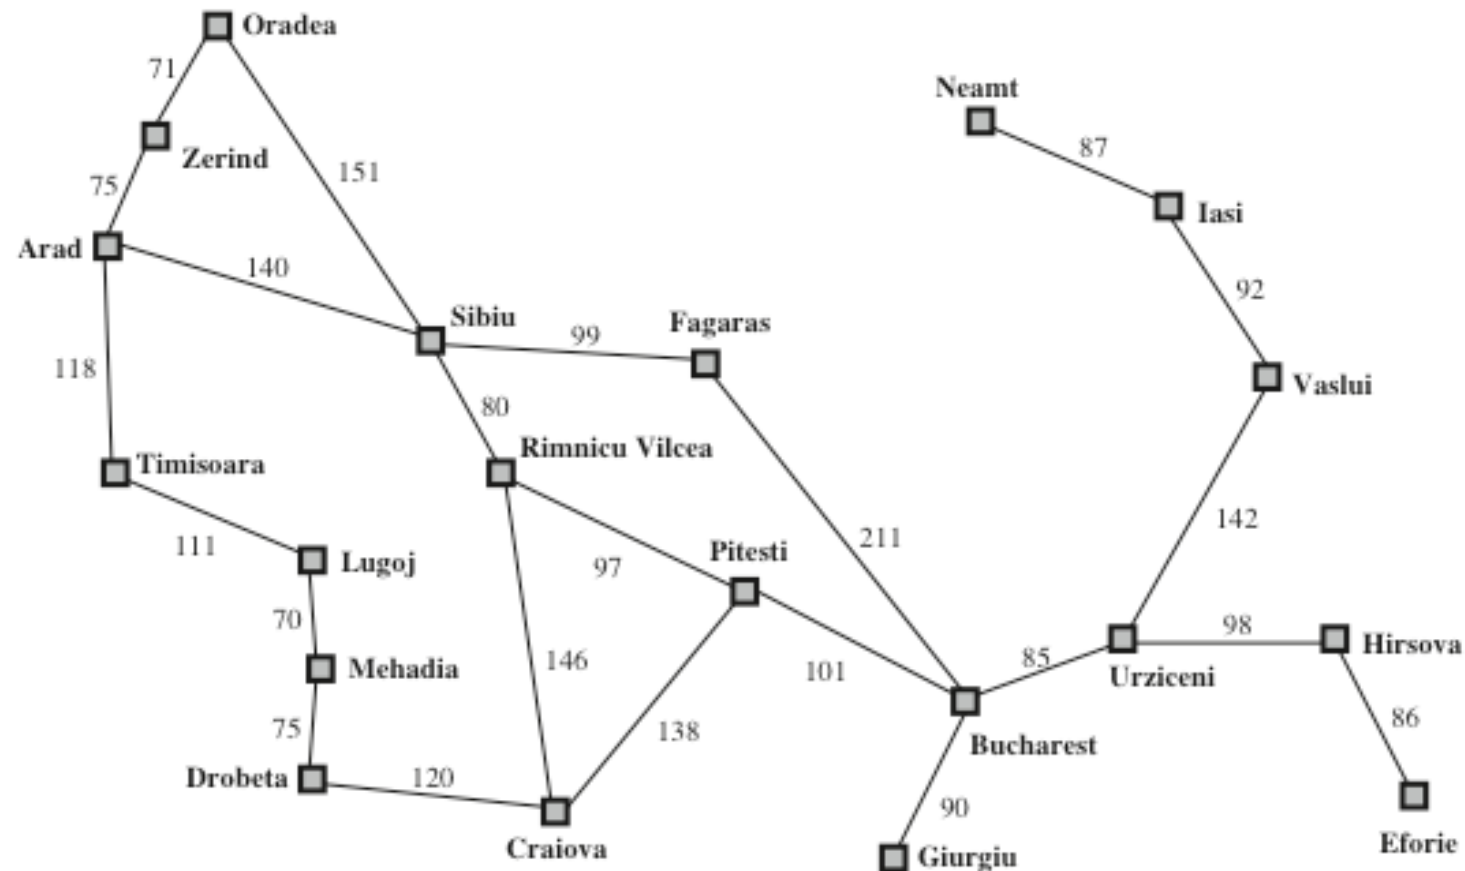
Properties of Greedy Best-First Search

15

- **Complete?** No – can get stuck in infinite loop, e.g., to get from Iasi to Fagaras, Iasi \rightarrow Neamt \rightarrow Iasi \rightarrow Neamt \rightarrow ...
- **Optimal?** No
- **Time?** $O(b^m)$, but a good heuristic can reduce complexity substantially
- **Space?** $O(b^m)$: keeps frontier nodes in memory

Romania with Step Costs (km)

16



A* Search

17

- Idea: Avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
- $g(n)$ = path cost from start node to reach n
- $h(n)$ = estimated cost of cheapest path from n to goal
- $f(n)$ = estimated cost of cheapest path through n to goal

A* Search Example

18

▶ Arad
366=0+366


A* Search Example

19



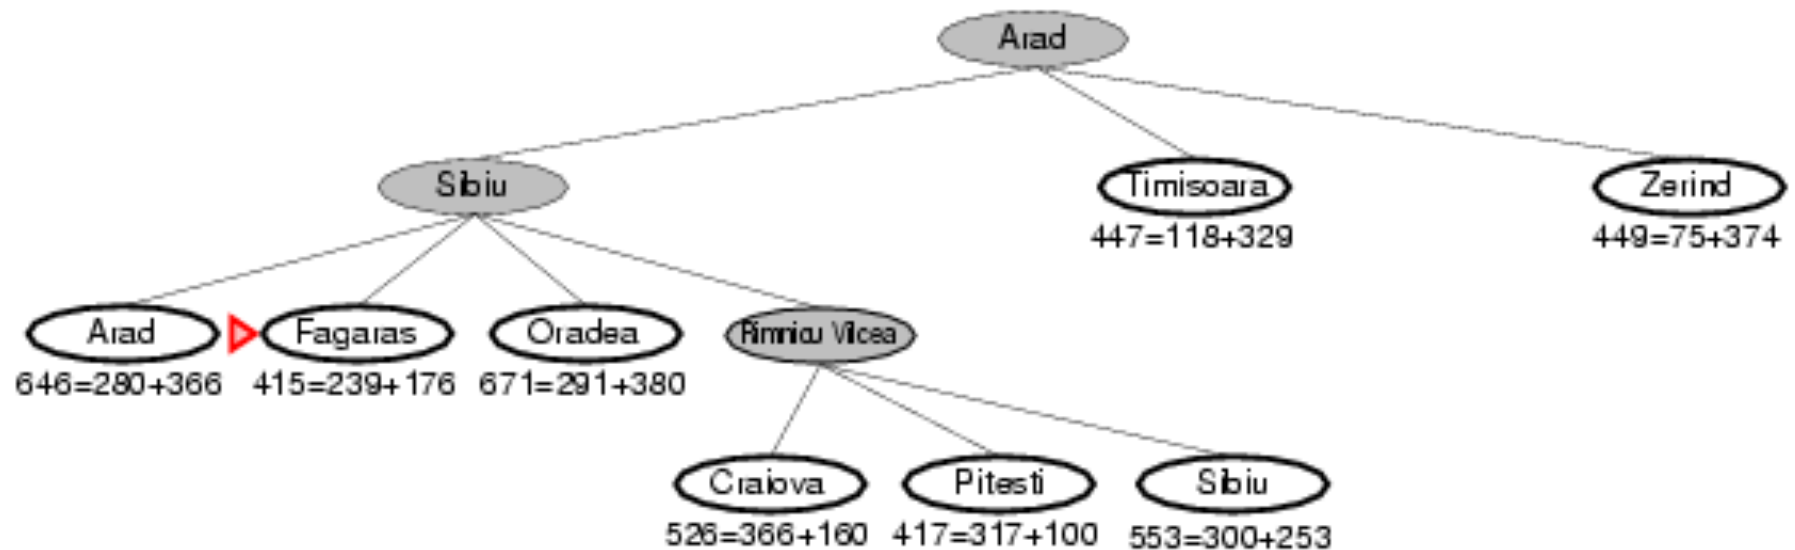
A* Search Example

20



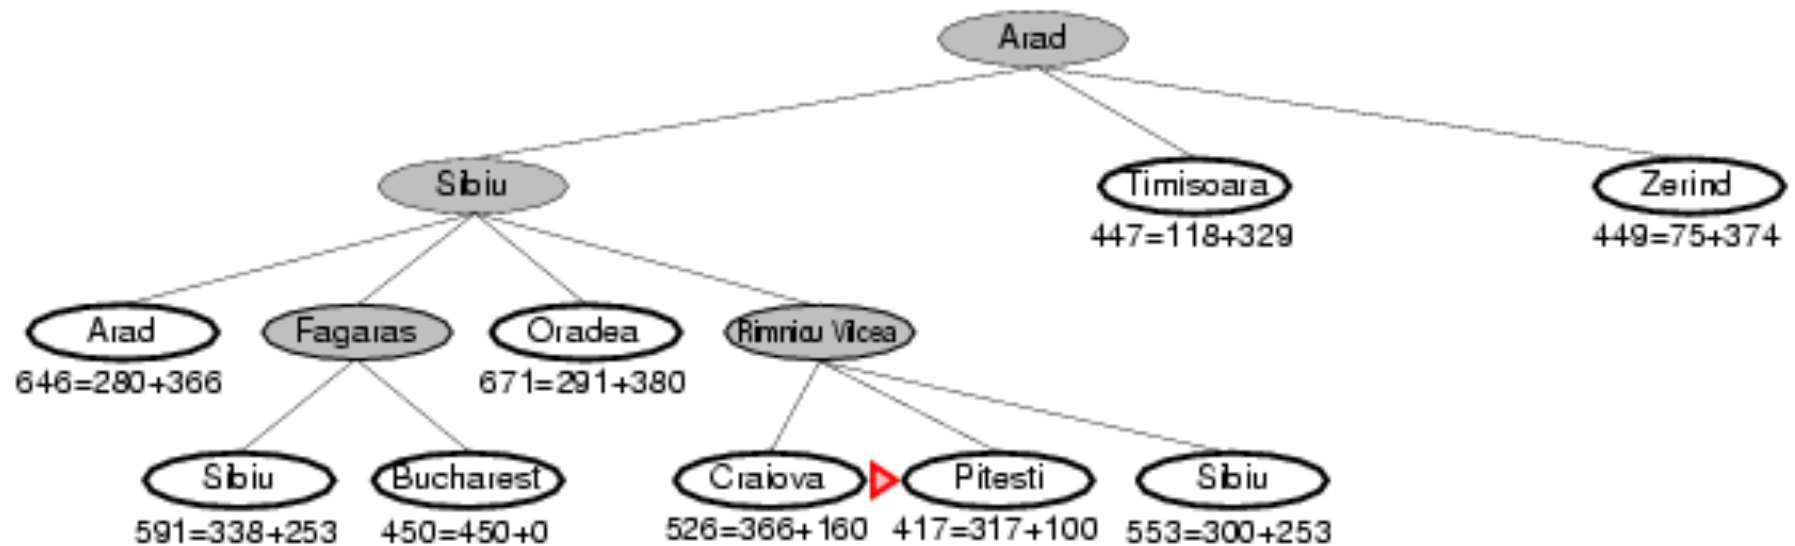
A* Search Example

21



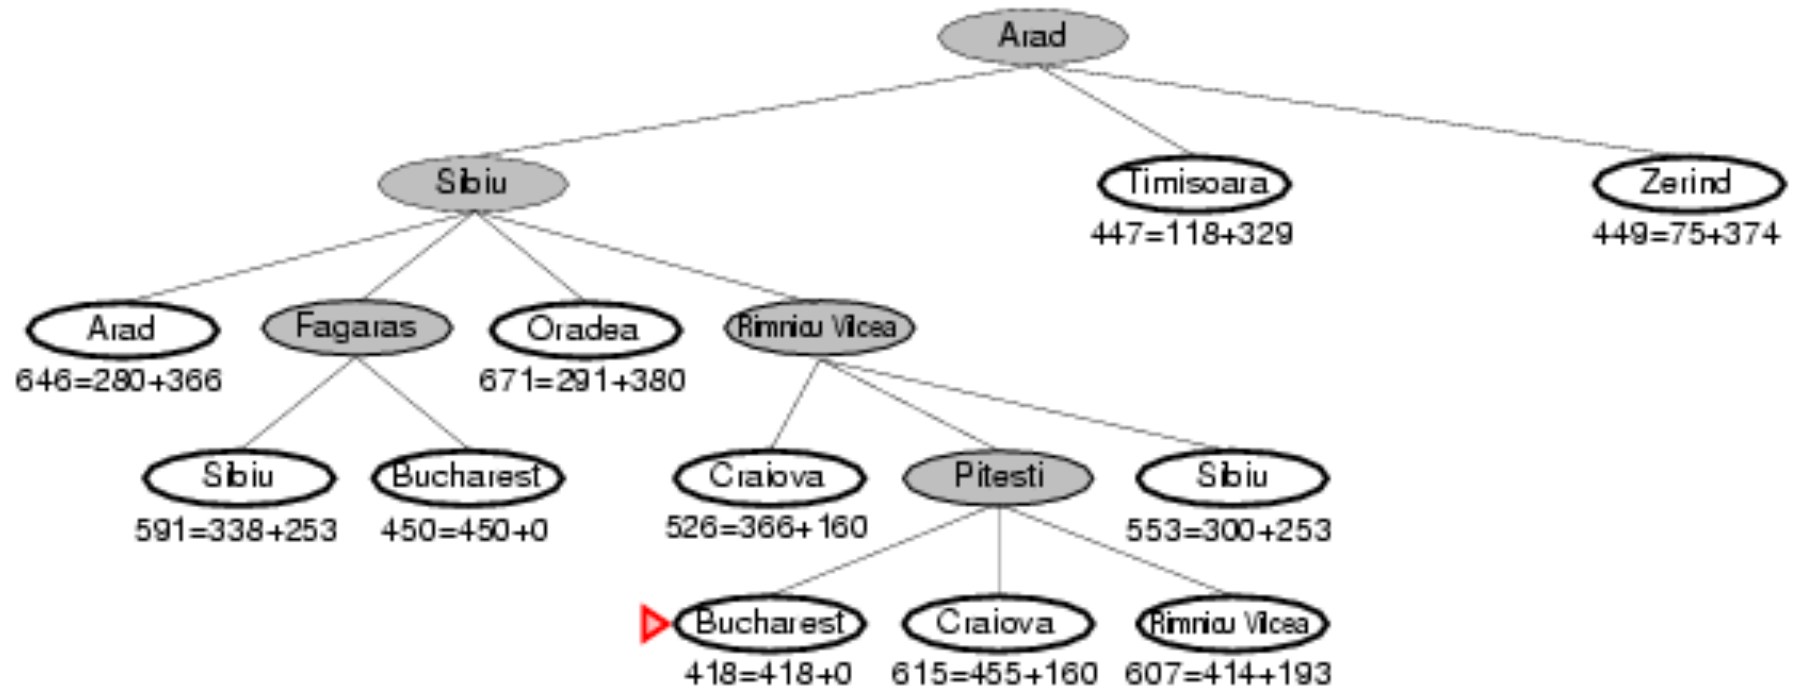
A* Search Example

22



A* Search Example

23



Admissible Heuristics

24


- A heuristic $h(n)$ is **admissible** if, for every node n , $h(n) \leq h^*(n)$ where $h^*(n)$ is the **true** cost to reach the goal state from n .
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**
- Example: $h_{SLD}(n)$ never overestimates the actual road distance
- **Theorem**: If $h(n)$ is admissible, then A* using TREE-SEARCH is optimal

Optimality of A^* using TREE-SEARCH

25

Proof Sketch:

- Suppose some suboptimal goal G_2 has been generated and is in the frontier. Let n be an unexpanded node in the frontier such that n is on a shortest path to an optimal goal G .




- $f(G) = g(G)$ since $h(G) = 0$
- $f(G_2) = g(G_2)$ since $h(G_2) = 0$
- $g(G_2) > g(G)$ since G_2 is suboptimal
- $f(G_2) > f(G)$ from above

Optimality of A^* using TREE-SEARCH

26

- Suppose some suboptimal goal G_2 has been generated and is in the frontier. Let n be an unexpanded node in the frontier such that n is on a shortest path to an optimal goal G .



- $f(G_2) > f(G)$ from previous slide
- $h(n) \leq h^*(n)$ since h is admissible
- $g(n) + h(n) \leq g(n) + h^*(n)$
- $f(n) \leq f(G)$
- Hence $f(G_2) > f(n)$, and A^* will never select G_2 for expansion

Consistent Heuristics


27

- A heuristic is **consistent** if, for every node n and every successor n' of n generated by any action a , $h(n) \leq c(n, a, n') + h(n')$.

- If h is consistent, we have

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n, a, n') + h(n') \\ &\geq g(n) + h(n) = f(n) \end{aligned}$$


- i.e., $f(n)$ is non-decreasing along any path.
- **Theorem:** If $h(n)$ is consistent, then A* using GRAPH-SEARCH is optimal



Optimality of A^* using GRAPH-SEARCH

28

- $f(n)$ is non-decreasing along any path
- A^* expands nodes in nondecreasing order of f value
- Gradually adds “ f -contours” of nodes
- Contour i has all nodes with $f = f_i$ where $f_i < f_{i+1}$



Optimality of A^* using GRAPH-SEARCH

29

Proof Sketch:

1. If $h(n)$ is consistent, then $f(n)$ is non-decreasing along any path
2. Whenever A^* selects a node n (e.g., repeated state) for expansion, the optimal path to n has been found
3. A^* expands nodes in nondecreasing order of $f(n)$
4. First expanded goal node is optimal solution:
 - a. f is true cost for goal nodes ($h(G) = 0$)
 - b. all later goal nodes are at least as expensive

Properties of A^*

30

- **Complete?** Yes (if there are finitely many nodes with $f(n) \leq f(G)$)
- **Optimal?** Yes
- **Time?** $O(b^{h^*-h})$ where h^* is actual cost of getting from root to goal
- **Space?** Keeps all generated nodes in memory

Admissible vs. Consistent Heuristics

31

- Why is consistency a stronger sufficient condition than admissibility? How are they related?
 - ▣ A consistent heuristic is admissible
 - ▣ An admissible heuristic MAY be inconsistent
- An admissible but inconsistent heuristic cannot guarantee optimality of A^* using GRAPH-SEARCH
 - ▣ Problem: GRAPH-SEARCH discards new paths to a repeated state. So, it may discard the optimal path. Previous proof breaks down.
 - ▣ Solution: Ensure that optimal path to any repeated state is always followed first.

Admissible Heuristics

32

- Let's revisit the 8-puzzle
 - ▣ Branching factor is about 3
 - ▣ Average solution depth is about 22 steps
 - ▣ Exhaustive tree search examines 3^{22} states
- How do we come up with good heuristics?

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Admissible Heuristics

33

E.g., 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance (i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$ 8
- $h_2(S) = ?$ $3+1+2+2+2+3+3+2 = 18$

Dominance

34

- If $h_2(n) \geq h_1(n)$ for all n (both admissible), then h_2 **dominates** h_1 .
It follows that h_2 incurs lower search cost than h_1 .

Average search costs (nodes generated):

- $d = 12$ IDS = 3,644,035 nodes
 $A^*(h_1) = 227$ nodes
 $A^*(h_2) = 73$ nodes
- $d = 24$ IDS = too many nodes
 $A^*(h_1) = 39,135$ nodes
 $A^*(h_2) = 1,641$ nodes

Deriving Admissible Heuristics

35

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem

Deriving Admissible Heuristics

36

- Rules of 8-puzzle:

A tile can move from square A to square B
if A is horizontally or vertically adjacent to B and B is blank

- We can generate three relaxed problems

1. A tile can move from square A to square B if A is adjacent to B
2. A tile can move from square A to square B if B is blank
3. A tile can move from square A to square B

Deriving Admissible Heuristics

37

- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ gives the shortest solution
- If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ gives the shortest solution

LOCAL SEARCH

ALMA Chapter 4.1

Outline

39

- Local search algorithms
 - ▣ Hill-climbing search
 - ▣ Simulated annealing
 - ▣ Local beam search
 - ▣ Genetic algorithms

Local Search Algorithms

40

- In many optimization problems, the **path** to goal is irrelevant; the goal state itself is the solution
- State space = set of “complete” configurations
- Find final configuration satisfying constraints, e.g., n-queens
- In such cases, we can use **local search algorithms** to keep a single “current” state and try to improve it
- Advantages: (1) very little/constant memory, and (2) find reasonable solutions in large state space

Example: n -queens

41

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



Hill-Climbing Search

42


- “Like climbing Mt. Everest in thick fog with amnesia”

```
function HILL-CLIMBING(problem) returns a state that is a local maximum  
  
  current  $\leftarrow$  MAKE-NODE(problem.INITIAL-STATE)  
  loop do  
    neighbor  $\leftarrow$  a highest-valued successor of current  
    if neighbor.VALUE  $\leq$  current.VALUE then return current.STATE  
    current  $\leftarrow$  neighbor
```

Hill-Climbing Search

43

- Problem: depending on initial state, can get stuck in local maxima



- Non-guaranteed fixes: sideway moves, random restarts

Hill-Climbing Search: 8-Queens


44

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♙	13	16	13	16
♙	14	17	15	♙	14	16	16
17	♙	16	18	15	♙	15	♙
18	14	♙	15	15	14	♙	16
14	14	13	17	12	14	12	18

- h = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = 17$ for the above state

Hill-Climbing Search: 8-Queens

45



A local minimum with $h = 1$

Simulated Annealing

46

- Idea: escape local maxima by allowing some “downhill” moves but **gradually decrease** their frequency

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”

  current  $\leftarrow$  MAKE-NODE(problem.INITIAL-STATE)
  for  $t = 1$  to  $\infty$  do
     $T \leftarrow$  schedule( $t$ )
    if  $T = 0$  then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  next.VALUE – current.VALUE
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 
```

Properties of Simulated Annealing


47

- Can prove: If T decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1
- Widely used in VLSI layout, factory scheduling, and other large-scale optimization tasks

Local Beam Search

48

- Why keep just one best state?



- Problem: May quickly concentrate in small region of state space → Solution: stochastic beam search

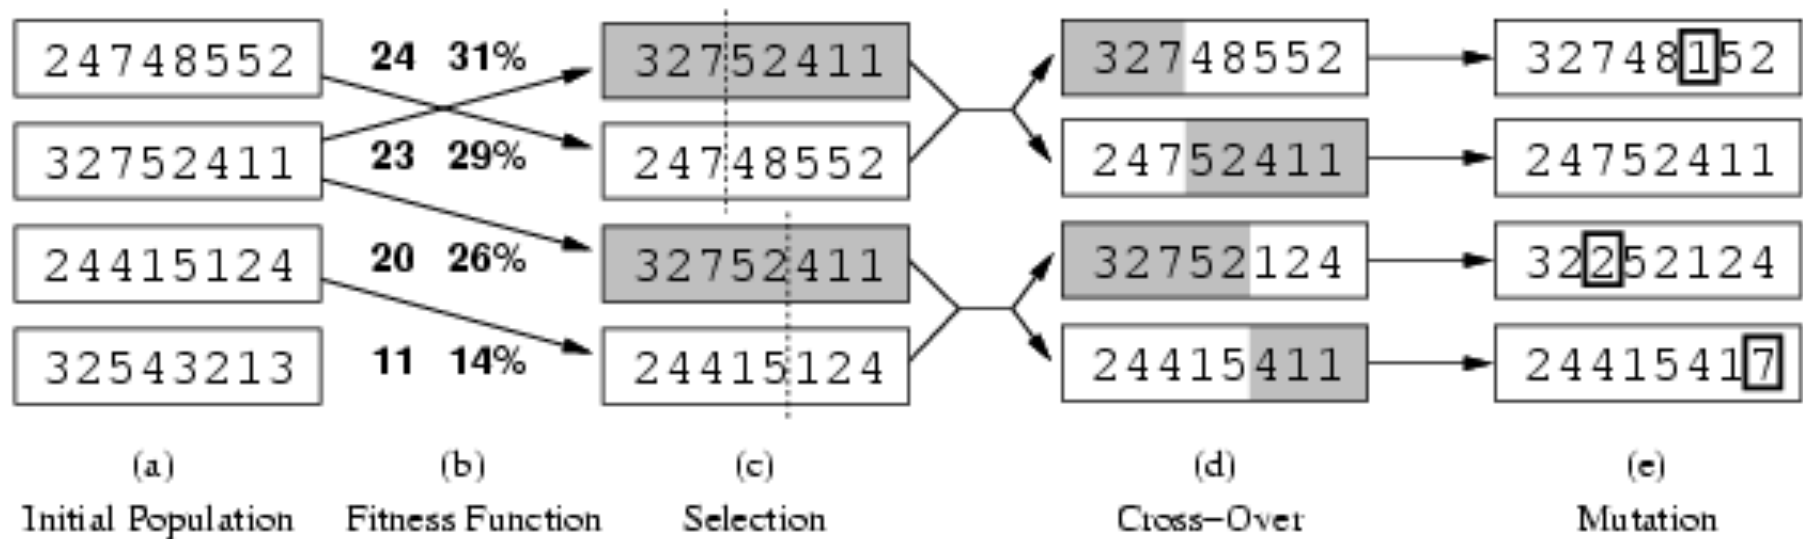
Genetic Algorithms

49

- Idea: a successor state is generated by combining two parent states
- Start with k randomly generated states (**population**)
- A state is represented as a string over a finite alphabet (often a string of 0s and 1s)
- Evaluation function (**fitness function**): higher values for better states
- Produce the next generation of states by selection, crossover, and mutation

Genetic Algorithms Example


50



- Fitness function: number of non-attacking pairs of queens (min = 0, max = $(8 \times 7)/2 = 28$)
- $24/(24+23+20+11) = 31\%$
- $23/(24+23+20+11) = 29\%$ etc.

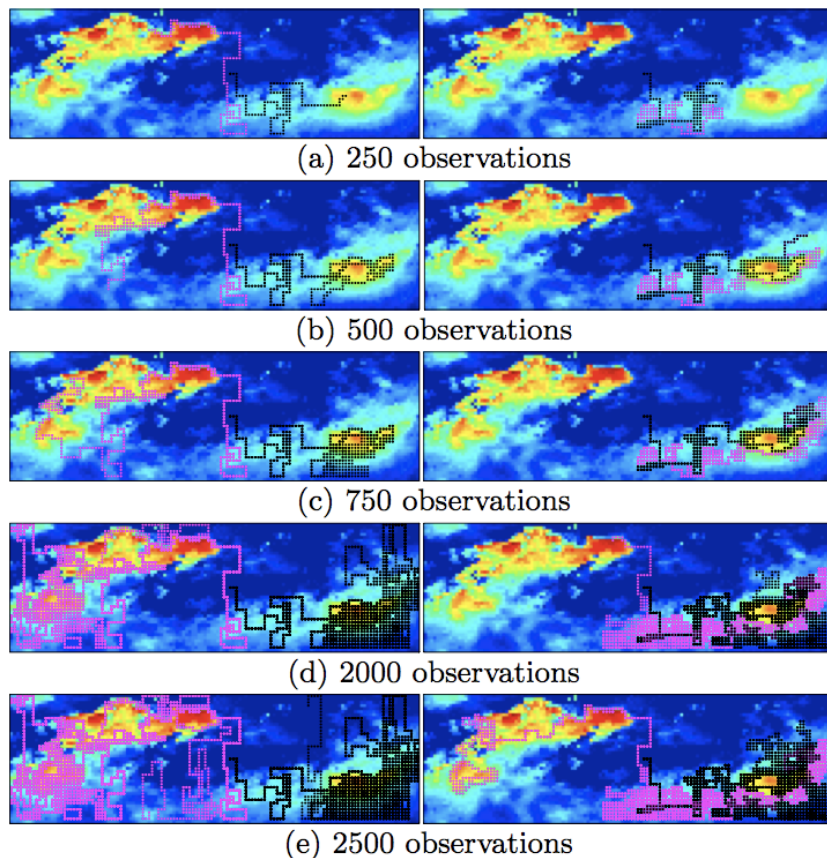
Genetic Algorithms: Crossover

51



Exploration Problems

52



Exploration problems: agent(s) physically in some part of the state space

- e.g. find hotspot using an agent or agents with local temperature sensors
- Sensible to expand states easily accessible to agent (i.e. **local** states)
 - Local search algorithms apply (e.g., hill-climbing)