

Previously...

1

- **Informed** search strategies use heuristics to guide search
 - ▣ Greedy best-first search
 - ▣ A* search
- If $h(n)$ is **admissible**, then A* using TREE-SEARCH is optimal
- If $h(n)$ is **consistent**, then A* using GRAPH-SEARCH is optimal
- A heuristic that **dominates** another incurs lower search cost

A* using GRAPH-SEARCH

2

UCS: $g(n)$ GBFS: $h(n)$ A*: $f(n) = g(n) + h(n)$

```
function UNIFORM-COST-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  frontier  $\leftarrow$  a priority queue ordered by PATH-COST, with node as the only element
  explored  $\leftarrow$  an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        frontier  $\leftarrow$  INSERT(child, frontier)
      else if child.STATE is in frontier with higher PATH-COST then
        replace that frontier node with child
```

Previously...

3

- **Local** search strategies
 - ▣ Hill-climbing search: use of heuristic function to improve “current” state
 - ▣ Simulated annealing: introduce some “downhill” moves to avoid getting trapped in local maxima
 - ▣ Local beam search: keep track of k best successor states
 - ▣ Genetic algorithms: mimic nature

ADVERSARIAL SEARCH

ALMA Chapter 5.1 – 5.5

Deep Blue vs. Garry Kasparov (1997)

5



Deterministic Games in Practice

6

- Checkers: **Chinook** ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a precomputed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- Chess: **Deep Blue** defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.
- Othello: human champions refuse to compete against computers, who are too good.
- Go: human champions refuse to compete against computers, who are too bad. In Go, $b > 300$, so most programs use pattern knowledge bases to suggest plausible moves.

Outline

7

- Adversarial search problems (aka games)
- Optimal (i.e., Minimax) decisions
- α - β pruning
- Imperfect, real-time decisions

Games vs. Search Problems


8

- “Unpredictable” opponent → solution is a strategy specifying a move for every possible opponent reply/response
- Time limit → unlikely to find goal, must approximate

Let's Play!

9

- Two players:



Game: Problem Formulation

10

A game is defined by 7 components:

1. Initial state S_0
2. States
3. Players : $\text{PLAYER}(s)$ defines which player has the move in state s
4. Actions : $\text{ACTIONS}(s)$ returns set of legal moves in state s
5. Transition model : $\text{RESULT}(s, a)$ returns state that results from the move a in state s

Game: Problem Formulation

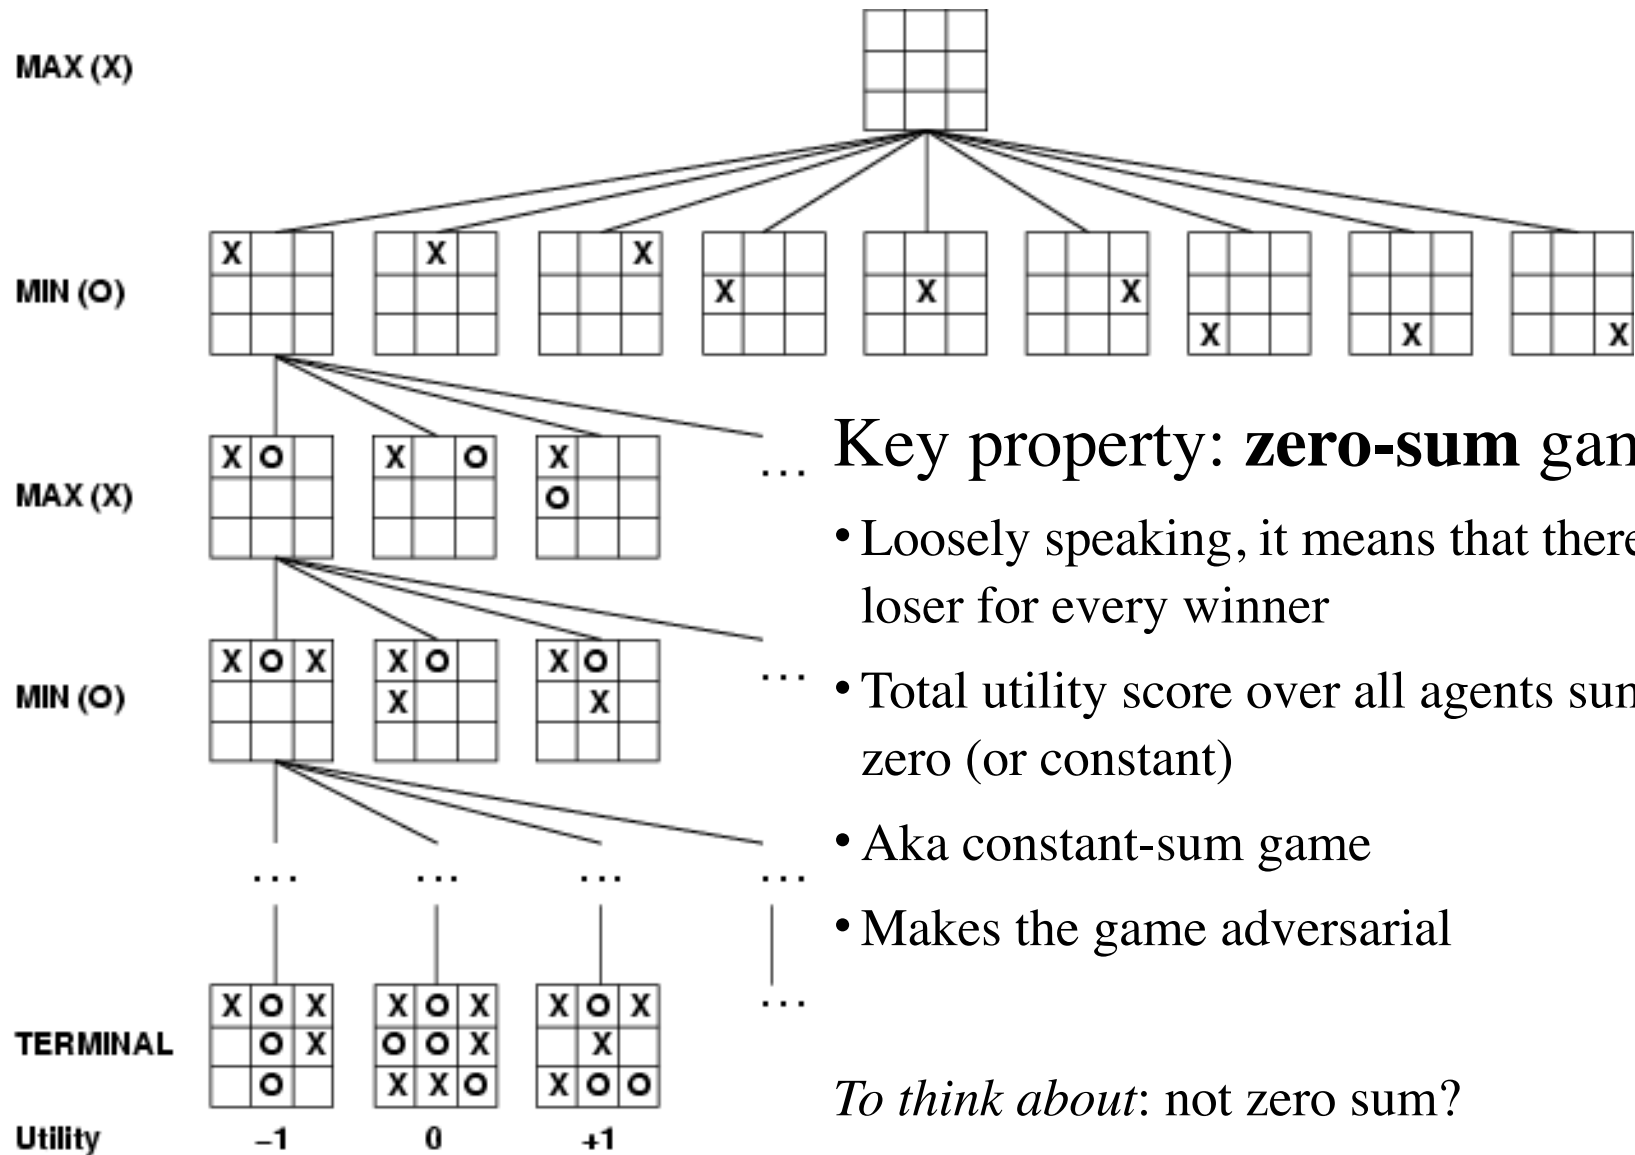
11

A game is defined by 7 components:

6. Terminal test $\text{TERMINAL-TEST}(s)$ returns true if game is over and false otherwise
 - Terminal states: states where the game has ended
7. Utility function $\text{UTILITY}(s, p)$ gives final numeric value for a game that ends in terminal state s for a player p
 - Example: for chess, win +1, loss -1, draw 0

Game Tree (2-Player, Deterministic, Turn-Taking)

12



Key property: **zero-sum** game

- Loosely speaking, it means that there's a loser for every winner
- Total utility score over all agents sum to zero (or constant)
- Aka constant-sum game
- Makes the game adversarial

To think about: not zero sum?

Example : Game of NIM

13

Several piles of sticks are given. We represent the configuration of the piles by a monotone sequence of integers, such as $(1,3,5)$. A player may remove, in one turn, any number of sticks from one pile. Thus, $(1,3,5)$ would become $(1,1,3)$ if the player were to remove 4 sticks from the last pile. The player who takes the last stick loses.


- Represent the NIM game $(1, 2, 2)$ as a game tree.



Minimax

14

- Optimal play for deterministic games
- Idea: choose move to state with highest **minimax value**
= best achievable payoff (for MAX) against MIN's optimal play
- E.g., 2-ply game:



Minimax Value

15

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Intuitively,

- MAX chooses move to maximize the minimum payoff
- MIN chooses move to minimize the maximum payoff

Minimax Algorithm

16

- Mistake in textbook (not reflected in errata): $s = state$

```
function MINIMAX-DECISION(state) returns an action  
  return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(state, a))$ 
```

```
function MAX-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$   
  return  $v$ 
```

```
function MIN-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$   
  return  $v$ 
```


Properties of Minimax

17


- **Complete?** Yes (if tree is finite)
- **Optimal?** Yes (against an optimal opponent)
- **Time complexity?** $O(b^m)$
- **Space complexity?** $O(bm)$ (depth-first exploration)
- For chess, $b \approx 35$, $m \approx 100$ for “reasonable” games \rightarrow time cost to find exact solution impractical
- Do we need to explore every path? **Pruning!**

α - β Pruning

18


- Idea: If we maintain the lower bound α and upper bound β of the values of, respectively, MAX's and MIN's nodes seen thus far, then we can prune subtrees that will never affect minimax decision.

19




α - β Pruning Example

20




α - β Pruning Example

21




α - β Pruning Example

22




α - β Pruning Example

23




α - β Pruning Example

24



α - β Pruning Example: Cannot Prune!

25



Properties of α - β Pruning

26

- Pruning **does not** affect final result
- Good move ordering improves effectiveness of pruning
- “Perfect” ordering: time complexity = $O(b^{m/2})$
 - ▣ **doubles** depth of search
 - ▣ E.g., for chess, simple ordering gets you close to this best-case result
 - ▣ Does it make sense then to have good heuristics for which nodes to expand first?
- Random ordering: time complexity = $O(b^{3m/4})$ for moderate b

Why is it called α - β Pruning?

27

- α = value of the best (i.e., highest-value) choice found so far at any choice point along the path for MAX
- If v is worse than α , then MAX will avoid it
→ prune remaining branches at node with v
- Define β similarly for MIN

MAX

MIN


..

..

..

MAX

MIN



The α - β Pruning Algorithm

28

function ALPHA-BETA-SEARCH(*state*) **returns** an action
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$
 return the *action* in ACTIONS(*state*) with value v

function MAX-VALUE(*state*, α , β) **returns** a utility value
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
 for each a **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \geq \beta$ **then return** v
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
 return v

function MIN-VALUE(*state*, α , β) **returns** a utility value
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow +\infty$
 for each a **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \leq \alpha$ **then return** v
 $\beta \leftarrow \text{MIN}(\beta, v)$
 return v

Summary: α - β Pruning Algorithm

29

- Initially, $\alpha = -\infty$, $\beta = +\infty$
- α is max along search path
- β is min along search path
- At MIN node, can stop if we find a node with value v smaller than or equal to α
- At MAX node, can stop if we find a node with value v larger than or equal to β

Time Limit

30

- Problem: very large search space in typical games
- Solution: α - β pruning removes large parts of search space
- Unresolved problem: Maximum depth of tree
- Standard solutions:
 - ▣ **evaluation function** = estimated expected utility of state
 - ▣ **cutoff test**: e.g., depth limit

Cutting Off Search

31

- Modify minimax or α - β pruning algorithms by replacing
 - ▣ `TERMINAL-TEST(state)` with `CUTOFF-TEST(state, depth)`
 - ▣ `UTILITY(state)` is replaced by `EVAL(state)`
- Can also be combined with iterative deepening

Heuristic Minimax Value

32

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

$$\text{H-MINIMAX}(s, d) = \begin{cases} \text{EVAL}(s) & \text{if } \text{CUTOFF-TEST}(s, d) \\ \max_{a \in \text{ACTIONS}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Designing Good Evaluation Functions

33

- Order all the terminal states in the same way as the true utility function
- Fast to compute
- For non-terminal states, must be strongly correlated with actual chances of winning

Evaluation Functions

34

- For chess, typically **linear** weighted sum of **features**

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$





- e.g., $w_1 = 9$ with

$f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$

- Caveat: Contribution of each feature is independent of other feature values
 - ▣ Bishops in chess better at endgame
- Should model the *expected utility value* over states with the same feature values

Expected Utility Value

35



- An expected utility value may map to many states, each of which may lead to different terminal states
- Want expected utility values to model likelihood of “better-utility” states
- Evaluation function need not return actual expected values as long as ordering of states is the same

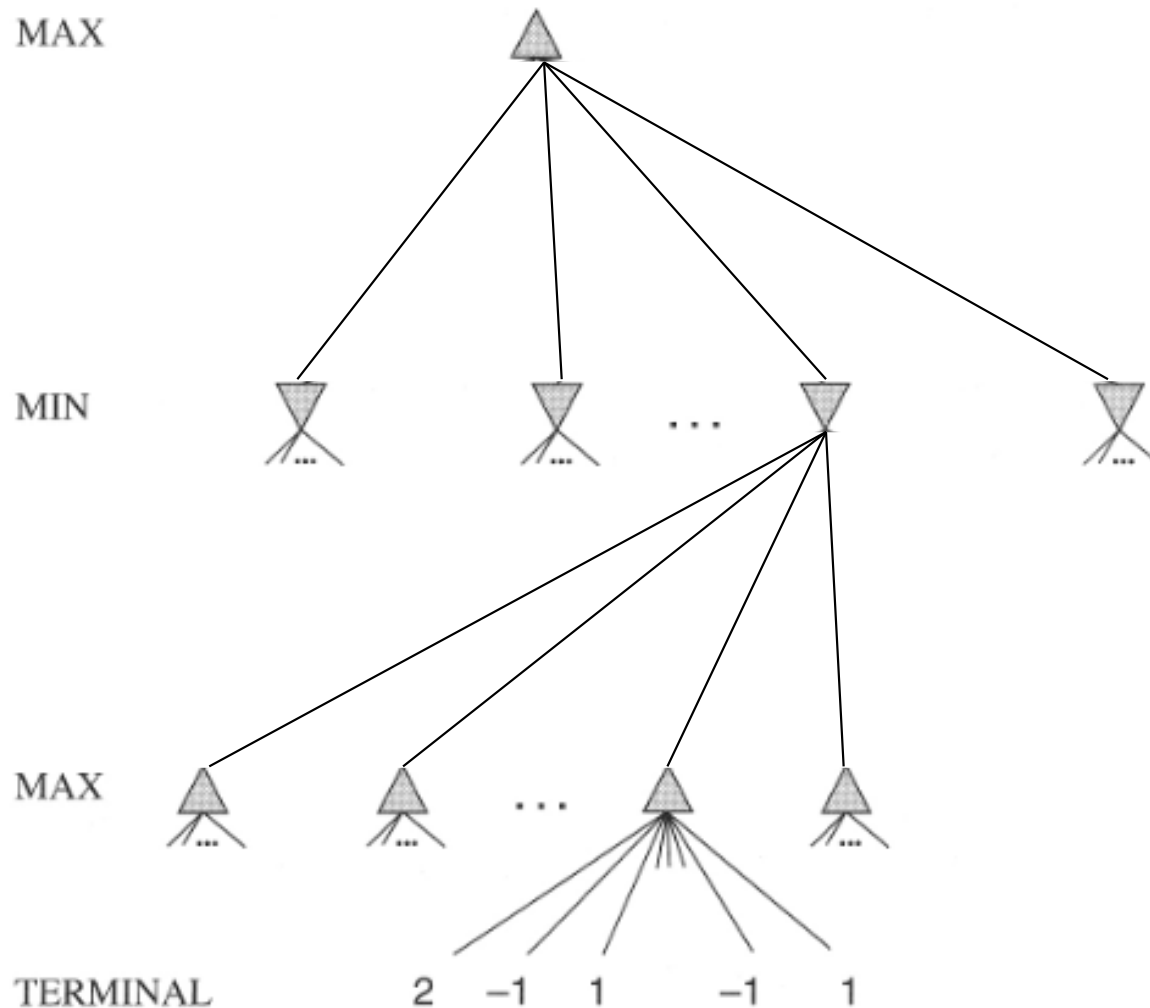
Stochastic Games

36

- What about an element of chance?
- How do we deal with uncertainty?
- Can we still use the minimax idea?

Adding Chance Layers

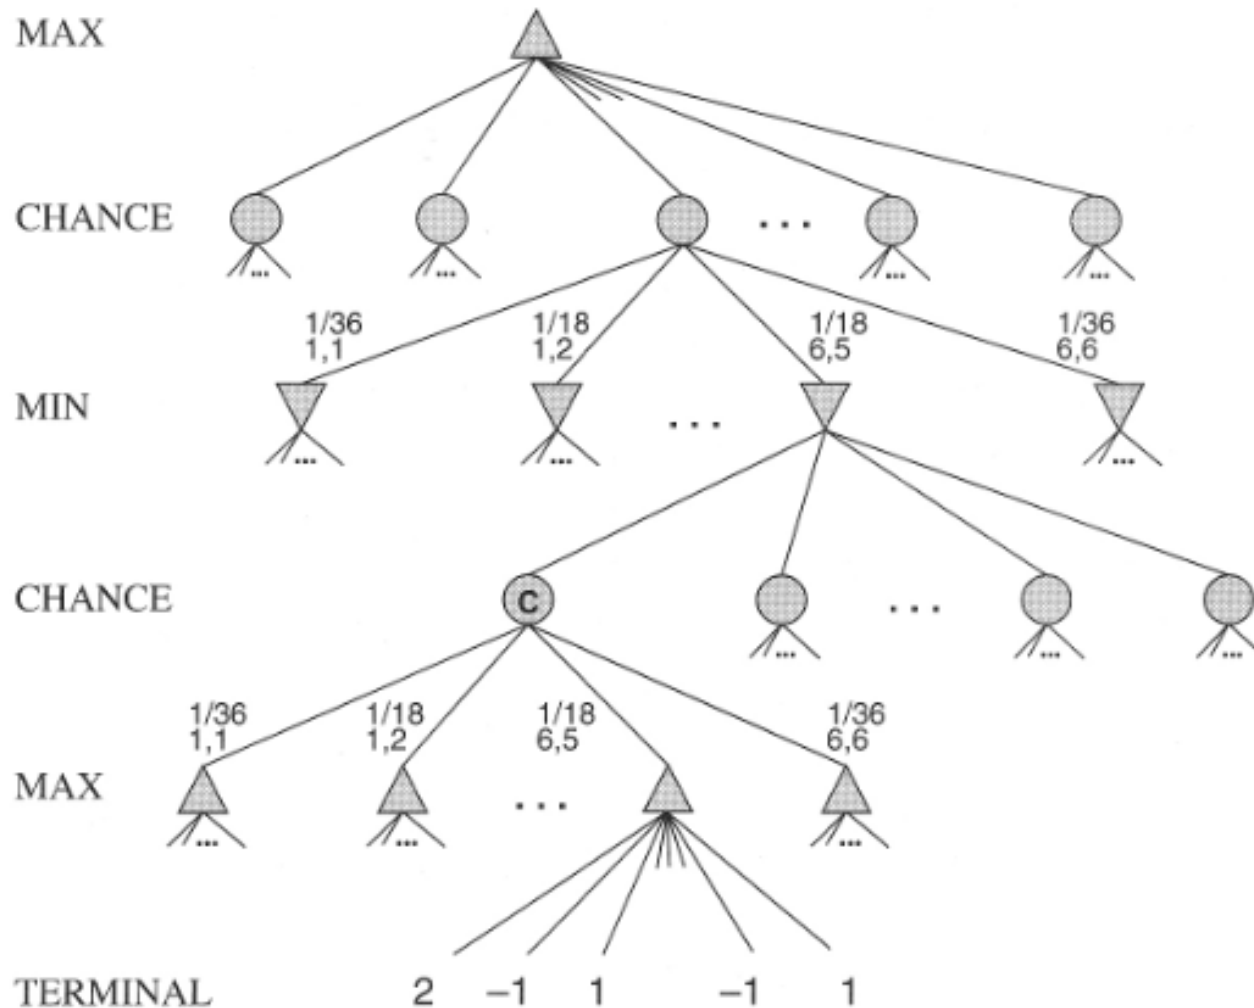
37



Calculate the **expected** value of a state

Adding Chance Layers

38



EXPECTIMINIMAX: Calculate the **expected** value of a state

Stochastic Games in Real World

39

